

清华大学本科生考试试题专用纸

考试课程：操作系统（A 卷）

时间：2013 年 06 月 22 日上午 8:00~10:00

系别：_____ 班级：_____ 学号：_____ 姓名：_____

答卷注意事项：1. 在开始答题前，请在试题纸和答卷本上写明系别、班级、学号和姓名。

2. 在答卷本上答题时，要写明题号，不必抄题。

3. 答题时，要书写清楚和整洁。

4. 请注意回答所有试题。本试卷有 6 个题目，共 20 页。

5. 考试完毕，必须将试题纸和答卷本一起交回。

一、(15 分) 在虚拟存储系统中，当由虚拟地址找不到对应的物理地址时，会产生缺页故障。请完成如下任务。

1) 描述缺页故障（page_fault）的处理流程；

2) 补全下面缺页处理中所缺代码。

=====kern/trap/trap.c=====

...

static int

pgfault_handler(struct trapframe *tf) {

extern struct mm_struct *check_mm_struct;

if(check_mm_struct !=NULL) { //used for test check_swap

 print_pgfault(tf);

 }

struct mm_struct *mm;

if (check_mm_struct != NULL) {

 assert(current == idleproc);

 mm = check_mm_struct;

 }

else {

if (current == NULL) {

 print_trapframe(tf);

 print_pgfault(tf);

 panic("unhandled page fault.\n");

 }

 mm = current->mm;

 }

return ... (1) ...;

}

...

```

static void
trap_dispatch(struct trapframe *tf) {
    char c;

    int ret=0;

    switch (tf->tf_trapno) {
    case T_DEBUG:
    case T_BRKPT:
        debug_monitor(tf);
        break;
    case T_PGFLT:
        if ((ret = ...(2)...) != 0) {
            print_trapframe(tf);
            if (current == NULL) {
                panic("handle pgfault failed. ret=%d\n", ret);
            }
            else {
                if (trap_in_kernel(tf)) {
                    panic("handle pgfault failed in kernel mode. ret=%d\n", ret);
                }
                cprintf("killed by kernel.\n");
                panic("handle user mode pgfault failed. ret=%d\n", ret);
                do_exit(-E_KILLED);
            }
        }
        break;
    case T_SYSCALL:
        syscall();
        break;
    case IRQ_OFFSET + IRQ_TIMER:
        // LAB3 : If some page replacement algorithm need tick to change the priority of pages,
        // then you can add code here.
        ticks ++;
        assert(current != NULL);
        run_timer_list();
        break;
    case IRQ_OFFSET + IRQ_COM1:
    case IRQ_OFFSET + IRQ_KBD:
        if ((c = cons_getc()) == 13) {
            debug_monitor(tf);
        }
        else {
            cprintf("%s [%03d] %c\n",
                    (tf->tf_trapno != IRQ_OFFSET + IRQ_KBD) ? "serial" : "kbd", c, c);

```

```

    }
    break;
case IRQ_OFFSET + IRQ_IDE1:
case IRQ_OFFSET + IRQ_IDE2:
    /* do nothing */
    break;
default:
    print_trapframe(tf);
    if (current != NULL) {
        cprintf("unhandled trap.\n");
        do_exit(-E_KILLED);
    }
    // in kernel, it must be a mistake
    panic("unexpected trap in kernel.\n");

}
}

=====kern/mm/vmm.c=====

// do_pgfault - interrupt handler to process the page fault execption
int
do_pgfault(struct mm_struct *mm, uint32_t error_code, uintptr_t addr) {
    int ret = -E_INVALID;
    struct vma_struct *vma = find_vma(mm, addr);

    pgfault_num++;
    if (vma == NULL ) {
        cprintf("not valid addr %x, and can not find it vma %x\n", addr, vma);
        goto failed;
    }
    else if (vma->vm_start > addr) {
        cprintf("not valid addr %x, and can not find it vma range[%x, %x]\n", addr, vma->vm_start
, vma->vm_end);
        goto failed;
    }
    cprintf("valid addr %x, and find it in vma range[%x, %x]\n", addr, vma->vm_start, vma->vm_end)
;

    switch (error_code & 3) {
default:
    /* default is 3: write, present */
case 2: /* write, not present */
    if (!(vma->vm_flags & VM_WRITE)) {
        cprintf("write, not present in do_pgfault failed\n");
        goto failed;
    }

```

```

    }
    break;
case 1: /* read, present */
    cprintf("read, present in do_pgfault failed\n");
    goto failed;
case 0: /* read, not present */
    if (!(vma->vm_flags & (VM_READ | VM_EXEC))) {
        cprintf("read, not present in do_pgfault failed\n");
        goto failed;
    }
}

uint32_t perm = PTE_U;
if (vma->vm_flags & VM_WRITE) {
    perm |= PTE_W;
}
addr = ROUNDDOWN(addr, PGSIZE);

ret = -E_NO_MEM;

pte_t *ptep;
// try to find a pte, if pte's PT(Page Table) isn't existed, then create a PT.
// (notice the 3th parameter '1')
if ((ptep = get_pte(mm->pgdir, addr, 1)) == NULL) {
    cprintf("get_pte in do_pgfault failed\n");
    goto failed;
}

if (*ptep == 0) { // if the phy addr isn't exist, then alloc a page & map the phy addr with logical addr
    if (...(3)... == NULL) {
        cprintf("pgdir_alloc_page in do_pgfault failed\n");
        goto failed;
    }
}
else {
    struct Page *page=NULL;
    cprintf("do pgfault: ptep %x, pte %x\n", ptep, *ptep);
    if (*ptep & PTE_P) {
        page = ...(4)...;
    } else{
        // if this pte is a swap entry, then load data from disk to a page with phy addr
        // and call page_insert to map the phy addr with logical addr
        if(swap_init_ok) {
            if ((ret = ...(5)...) != 0) {
                cprintf("swap_in in do_pgfault failed\n");
            }
        }
    }
}

```

```

        goto failed;
    }

}

else {
    cprintf("no swap_init_ok but ptep is %x, failed\n",*ptep);
    goto failed;
}

}

page_insert(mm->pgdir, page, addr, perm);
swap_map_swappable(mm, addr, page, 1);
}

ret = 0;
failed:
    return ret;
}

=====kern/mm/swap.c=====

...
int
swap_out(struct mm_struct *mm, int n, int in_tick)
{
    int i;
    for (i = 0; i != n; ++ i)
    {
        uintptr_t v;
        //struct Page **ptr_page=NULL;
        struct Page *page;
        // cprintf("i %d, SWAP: call swap_out_victim\n",i);
        int r = sm->swap_out_victim(mm, &page, in_tick);
        if (r != 0) {
            cprintf("i %d, swap_out: call swap_out_victim failed\n",i);
            break;
        }
        //assert(!PageReserved(page));

        //cprintf("SWAP: choose victim page 0x%08x\n", page);

        v=page->pra_vaddr;
        pte_t *ptep = get_pte(mm->pgdir, v, 0);
        assert((*ptep & PTE_P) != 0);

        if (...(6)... != 0) {

```

```

        cprintf("SWAP: failed to save\n");
        sm->map_swappable(mm, v, page, 0);
        continue;
    }
    else {
        cprintf("swap_out: i %d, store page in vaddr 0x%x to disk swap entry %d\n", i,
v, page->pra_vaddr/PGSIZE+1);
        *ptep = (page->pra_vaddr/PGSIZE+1)<<8;
        free_page(page);
    }

    tlb_invalidate(mm->pgdir, v);
}
return i;
}
int
swap_in(struct mm_struct *mm, uintptr_t addr, struct Page **ptr_result)
{
    struct Page *result = alloc_page();
    assert(result!=NULL);

    pte_t *ptep = get_pte(mm->pgdir, addr, 0);
    // cprintf("SWAP: load ptep %x swap entry %d to vaddr 0x%08x, page %x, No %d\n", ptep, (*ptep
)>>8, addr, result, (result->pages));

    int r;
    if ((r = ... (7) ...) != 0)
    {
        assert(r!=0);
    }
    cprintf("swap_in: load disk swap entry %d with swap_page in vadr 0x%x free_area.nr_free %d\n"
, (*ptep)>>8, addr, free_area.nr_free);
    *ptr_result=result;
    return 0;
}

=====kern/mm/pmm.h=====

...
#define alloc_page() alloc_pages(1)
#define free_page(page) free_pages(page, 1)

=====kern/mm/pmm.c=====

```

```

...

// pgdir_alloc_page - call alloc_page & page_insert functions to
//                     - allocate a page size memory & setup an addr map
//                     - pa<->la with linear address la and the PDT pgdir
struct Page *
pgdir_alloc_page(pde_t *pgdir, uintptr_t la, uint32_t perm) {
    struct Page *page = alloc_page();
    if (page != NULL) {
        if (page_insert(pgdir, page, la, perm) != 0) {
            free_page(page);
            return NULL;
        }
        if (swap_init_ok){
            if(check_mm_struct!=NULL) {
                swap_map_swappable(check_mm_struct, la, page, 0);
                page->pra_vaddr=la;
                assert(page_ref(page) == 1);
                //cprintf("get No. %d page: pra_vaddr %x, pra_link.prev %x, pra_link_next %x in p
gdir_alloc_page\n", (page-pages), page->pra_vaddr,page->pra_page_link.prev, page->pra_page_link.ne
xt);
            }
            else { //now current is existed, should fix it in the future
                //swap_map_swappable(current->mm, la, page, 0);
                //page->pra_vaddr=la;
                //assert(page_ref(page) == 1);
                //panic("pgdir_alloc_page: no pages. now current is existed, should fix it in the
future\n");
            }
        }
    }

    return page;
}

=====kern/fs/swapfs.c=====

...

int
swapfs_read(swap_entry_t entry, struct Page *page) {
    return ide_read_secs(SWAP_DEV_NO, swap_offset(entry) * PAGE_NSECT, page2kva(page), PAGE_NSECT)
;
}

```

```

int
swapfs_write(swap_entry_t entry, struct Page *page) {
    return ide_write_secs(SWAP_DEV_NO, swap_offset(entry) * PAGE_NSECT, page2kva(page), PAGE_NSECT);
};
}

```

=====kern/mm/swap_fifo.c=====

```

...
struct swap_manager swap_manager_fifo =
{
    .name          = "fifo swap manager",
    .init          = &_amp;_fifo_init,
    .init_mm       = &_amp;_fifo_init_mm,
    .tick_event    = &_amp;_fifo_tick_event,
    .map_swappable = &_amp;_fifo_map_swappable,
    .set_unswappable = &_amp;_fifo_set_unswappable,
    .swap_out_victim = &_amp;_fifo_swap_out_victim,
    .check_swap    = &_amp;_fifo_check_swap,
};

```

二、(15 分)调度器是操作系统内核中依据调度算法进行进程切换选择的模块。请完成如下任务。

- 1) 试描述步进调度算法(Stride Scheduling)的基本原理。
- 2) 请给出下面测试程序 (user/priority.c) 执行时的进程调度顺序。建议说明每次进程切换后当前执行进程的 ID、lab6_priority、lab6_stride 和已切换次数。

=====kern/process/proc.h=====

```

...
struct proc_struct {
    enum proc_state state;           // Process state
    int pid;                         // Process ID
    int runs;                        // the running times of Proces
    uintptr_t kstack;               // Process kernel stack
    volatile bool need_resched;      // bool value: need to be rescheduled to release C
PU?
    struct proc_struct *parent;      // the parent process
    struct mm_struct *mm;            // Process's memory management field
    struct context context;          // Switch here to run process
    struct trapframe *tf;            // Trap frame for current interrupt
    uintptr_t cr3;                   // CR3 register: the base addr of Page Directroy T
able(PDT)

```



```

uint32_t flags;                                // Process flag
char name[PROC_NAME_LEN + 1];                 // Process name
list_entry_t list_link;                       // Process link list
list_entry_t hash_link;                      // Process hash list
int exit_code;                               // exit code (be sent to parent proc)
uint32_t wait_state;                         // waiting state
struct proc_struct *cptr, *yptr, *optr;      // relations between processes
struct run_queue *rq;                        // running queue contains Process
list_entry_t run_link;                       // the entry linked in run queue
int time_slice;                             // time slice for occupying the CPU
skew_heap_entry_t lab6_run_pool;             // FOR LAB6 ONLY: the entry in the run pool
uint32_t lab6_stride;                        // FOR LAB6 ONLY: the current stride of the proces
s
uint32_t lab6_priority;                      // FOR LAB6 ONLY: the priority of process, set by
lab6_set_priority(uint32_t)
};

```

=====user/priority.c=====

```

#include <ulib.h>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

#define TOTAL 6
/* to get enough accuracy, MAX_TIME (the running time of each process) should >1000 mseconds. */
#define MAX_TIME 2000
unsigned int acc[TOTAL];
int status[TOTAL];
int pids[TOTAL];

static void
spin_delay(void)
{
    int i;
    volatile int j;
    for (i = 0; i != 200; ++ i)
    {
        j = !j;
    }
}

```

```

int
main(void) {
    int i,time;
    memset(pids, 0, sizeof(pids));
    lab6_set_priority(TOTAL + 1);

    for (i = 0; i < TOTAL; i ++) {
        acc[i]=0;
        if ((pids[i] = fork()) == 0) {
            lab6_set_priority(i + 1);
            acc[i] = 0;
            while (1) {
                spin_delay();
                ++ acc[i];
                if(acc[i]%4000==0) {
                    if((time=gettime_msec())>MAX_TIME) {
                        cprintf("child pid %d, acc %d, time %d\n",getpid(),acc[i],time);
                        exit(acc[i]);
                    }
                }
            }
        }
        if (pids[i] < 0) {
            goto failed;
        }
    }

    cprintf("main: fork ok,now need to wait pids.\n");

    for (i = 0; i < TOTAL; i ++) {
        status[i]=0;
        waitpid(pids[i],&status[i]);
        cprintf("main: pid %d, acc %d, time %d\n",pids[i],status[i],gettime_msec());
    }
    cprintf("main: wait pids over\n");
    cprintf("stride sched correct result:");
    for (i = 0; i < TOTAL; i ++)
    {
        cprintf(" %d", (status[i] * 2 / status[0] + 1) / 2);
    }
    cprintf("\n");

    return 0;

```


failed:

```

    for (i = 0; i < TOTAL; i++) {
        if (pids[i] > 0) {
            kill(pids[i]);
        }
    }
    panic("FAIL: T.T\n");
}

```

三、(20 分)一个从键盘输入到打印机输出的数据处理流程图如下图所示。其中键盘输入进程 (input_process) 通过缓冲区 Buffer1 把数据传送给计算进程 (calculation_process)，计算进程把处理结果通过 Buffer2 传送给打印进程 (output_process)。键盘输入进程、计算进程及打印进程对缓冲区 Buffer1 和 Buffer2 的访问满足下列条件。



- 1) 任何时刻只有一个进程在对缓冲区 Buffer1 进行数据读写操作；只有一个进程在对缓冲区 Buffer2 进行数据读写操作；允许进程对 Buffer1 和 Buffer2 的同时读写操作。
- 2) 两个缓冲区的大小是无限大。

请用信号量 (semaphore) 机制实现键盘输入进程 input_process()、计算进程 calculation_process() 和打印进程 output_process()。要求：用信号量方法 (不允许使用信号量集)，并给出信号量定义和初始值；在代码中要有适当的注释，以说明信号量定义的作用和代码的含义；用类 C 语言描述共享变量和函数。

四、Bakery 算法(Lamport 1979)是一种解决 n 个线程访问临界区 (Critical Section) 问题的软件同步算法。该算法定义了两个共享数组如下：

```
boolean choosing[n];
```

```
int number[n]; //ticket
```

所有 choosing[i] 的初始值为 false，而 number[i] 的初始值为 0。进程 i 访问临界区的伪代码如下。

```

do {
    choosing[i] = true;
    number[i] = max(number[0], number[1], ..., number [n - 1])+1;
    choosing[i] = false;

    for (j = 0; j < n; j++) {
        while (choosing[j]) ; // (A)
        while ((number[j] != 0) && ( (number[j],j) < (number[i],i) ) ) ;
    }

    critical section
    number[i] = 0;
    remainder section
} while (1);

```

请完成如下任务。

- 1) 删除 (A) 行的代码, 算法还能正确控制对临界区的访问吗?
- 2) 如果能, 请说明理由; 如果不能, 请给出一个出现错误的情形, 并解释原因。

五、(20 分) 在 ucore 中采用的文件系统是 UNIX 文件系统 UFS 的简化版本 SFS。请完成如下任务。

- 1) 描述 UFS 的多级间接索引文件 (Multi-level Indexed Allocation) 的存储结构;
- 2) 补全下面文件系统代码。

=====kern/fs/sfs/sfs.h=====

```
#ifndef __KERN_FS_SFS_SFS_H__
#define __KERN_FS_SFS_SFS_H__

#include <defs.h>
#include <mmu.h>
#include <list.h>
#include <sem.h>
#include <unistd.h>

#define SFS_MAGIC                0x2f8dbe2a                /* magic number for sfs */
#define SFS_BLKSIZE              PGSIZE                    /* size of block */
#define SFS_NDIRECT              12                        /* # of direct blocks in inode */
#define SFS_MAX_INFO_LEN        31                        /* max length of information */
#define SFS_MAX_FNAME_LEN        FS_MAX_FNAME_LEN          /* max length of filename */
#define SFS_MAX_FILE_SIZE        (1024UL * 1024 * 128)      /* max file size (128M) */
#define SFS_BLKKN_SUPER          0                          /* block the superblock lives in */
#define SFS_BLKKN_ROOT           1                          /* location of the root dir inode */
#define SFS_BLKKN_FREEMAP        2                          /* 1st block of the freemap */

/* # of bits in a block */
#define SFS_BLKBITS               (SFS_BLKSIZE * CHAR_BIT)

/* # of entries in a block */
```

```

#define SFS_BLK_NENTRY                                (SFS_BLKSIZE / sizeof(uint32_t))

/* file types */
#define SFS_TYPE_INVAL                                0      /* Should not appear on disk */
#define SFS_TYPE_FILE                                  1
#define SFS_TYPE_DIR                                  2
#define SFS_TYPE_LINK                                  3

/*
 * On-disk superblock
 */
struct sfs_super {
    uint32_t magic;                                /* magic number, should be SFS_MAGIC */
    uint32_t blocks;                                /* # of blocks in fs */
    uint32_t unused_blocks;                          /* # of unused blocks in fs */
    char info[SFS_MAX_INFO_LEN + 1];                /* infomation for sfs */
};

/* inode (on disk) */
struct sfs_disk_inode {
    uint32_t size;                                /* size of the file (in bytes) */
    uint16_t type;                                /* one of SYS_TYPE_* above */
    uint16_t nlinks;                                /* # of hard links to this file */
    uint32_t blocks;                                /* # of blocks */
    uint32_t direct[SFS_NDIRECT];                  /* direct blocks */
    uint32_t indirect;                              /* indirect blocks */
    //    uint32_t db_indirect;                        /* double indirect blocks */
    //    unused
};

/* file entry (on disk) */
struct sfs_disk_entry {
    uint32_t ino;                                /* inode number */
    char name[SFS_MAX_FNAME_LEN + 1];                /* file name */
};

#define sfs_dentry_size                                \
    sizeof(((struct sfs_disk_entry *)0)->name)

/* inode for sfs */
struct sfs_inode {
    struct sfs_disk_inode *din;                    /* on-disk inode */
    uint32_t ino;                                /* inode number */
    bool dirty;                                    /* true if inode modified */
    int reclaim_count;                              /* kill inode if it hits zero */
    semaphore_t sem;                                /* semaphore for din */
};

```

```

    list_entry_t inode_link;                /* entry for linked-list in sfs_fs */
    list_entry_t hash_link;                /* entry for hash linked-list in sfs_fs */
};

#define le2sin(le, member)                \
    to_struct((le), struct sfs_inode, member)

/* filesystem for sfs */
struct sfs_fs {
    struct sfs_super super;                /* on-disk superblock */
    struct device *dev;                    /* device mounted on */
    struct bitmap *freemap;                /* blocks in use are mared 0 */
    bool super_dirty;                    /* true if super/freemap modified */
    void *sfs_buffer;                    /* buffer for non-block aligned io */
    semaphore_t fs_sem;                    /* semaphore for fs */
    semaphore_t io_sem;                    /* semaphore for io */
    semaphore_t mutex_sem;                /* semaphore for link/unlink and rename */
    list_entry_t inode_list;                /* inode linked-list */
    list_entry_t *hash_list;                /* inode hash linked-list */
};

/* hash for sfs */
#define SFS_HLIST_SHIFT                    10
#define SFS_HLIST_SIZE                    (1 << SFS_HLIST_SHIFT)
#define sin_hashfn(x)                    (hash32(x, SFS_HLIST_SHIFT))

/* size of freemap (in bits) */
#define sfs_freemap_bits(super)            ROUNDUP((super)->blocks, SFS_BLKBITS)

/* size of freemap (in blocks) */
#define sfs_freemap_blocks(super)        ROUNDUP_DIV((super)->blocks, SFS_BLKBITS)

struct fs;
struct inode;

void sfs_init(void);
int sfs_mount(const char *devname);

void lock_sfs_fs(struct sfs_fs *sfs);
void lock_sfs_io(struct sfs_fs *sfs);
void lock_sfs_mutex(struct sfs_fs *sfs);
void unlock_sfs_fs(struct sfs_fs *sfs);
void unlock_sfs_io(struct sfs_fs *sfs);
void unlock_sfs_mutex(struct sfs_fs *sfs);

int sfs_rblock(struct sfs_fs *sfs, void *buf, uint32_t blkno, uint32_t nblks);

```

```

int sfs_wblock(struct sfs_fs *sfs, void *buf, uint32_t blkno, uint32_t nblks);
int sfs_rbuf(struct sfs_fs *sfs, void *buf, size_t len, uint32_t blkno, off_t offset);
int sfs_wbuf(struct sfs_fs *sfs, void *buf, size_t len, uint32_t blkno, off_t offset);
int sfs_sync_super(struct sfs_fs *sfs);
int sfs_sync_freemap(struct sfs_fs *sfs);
int sfs_clear_block(struct sfs_fs *sfs, uint32_t blkno, uint32_t nblks);

int sfs_load_inode(struct sfs_fs *sfs, struct inode **node_store, uint32_t ino);

#endif /* !__KERN_FS_SFS_SFS_H__ */

=====kern/fs/sfs/sfs_inode.c=====

.....
static int
sfs_bmap_get_sub_nolock(struct sfs_fs *sfs, uint32_t *entp, uint32_t index, bool create, uint32_t
*ino_store) {
    assert(index < SFS_BLK_NENTRY);
    int ret;
    uint32_t ent, ino = 0;
    off_t offset = index * sizeof(uint32_t);
    if ((ent = *entp) != 0) {
        if ((ret = ...(1)... ) != 0) {
            return ret;
        }
        if (ino != 0 || !create) {
            goto out;
        }
    }
    else {
        if (!create) {
            goto out;
        }
        if ((ret = sfs_block_alloc(sfs, &ent)) != 0) {
            return ret;
        }
    }

    if ((ret = sfs_block_alloc(sfs, &ino)) != 0) {
        goto failed_cleanup;
    }
    if ((ret = ...(2)... ) != 0) {
        sfs_block_free(sfs, ino);
        goto failed_cleanup;
    }

```

```

    }

out:
    if (ent != *entp) {
        *entp = ent;
    }
    *ino_store = ino;
    return 0;

failed_cleanup:
    if (ent != *entp) {
        sfs_block_free(sfs, ent);
    }
    return ret;
}

.....
static int
sfs_bmap_get_nolock(struct sfs_fs *sfs, struct sfs_inode *sin, uint32_t index, bool create, uint32
_t *ino_store) {
    struct sfs_disk_inode *din = sin->din;
    int ret;
    uint32_t ent, ino;
    if (...(3)...) {
        if ((ino = din->direct[index]) == 0 && create) {
            if ((ret = sfs_block_alloc(sfs, &ino)) != 0) {
                return ret;
            }
            din->direct[index] = ino;
            sin->dirty = 1;
        }
        goto out;
    }

    index = ...(4)...;
    if (index < SFS_BLK_NENTRY) {
        ent = din->indirect;
        if ((ret = ...(5)...) != 0) {
            return ret;
        }
        if (ent != din->indirect) {
            assert(din->indirect == 0);
            din->indirect = ent;
            sin->dirty = 1;
        }
        goto out;
    }

```



```

    }

    index = ...(6)...;
    if ((ent = ino) != 0) {
        if ((ret = sfs_bmap_get_sub_nolock(sfs, &ent, index % SFS_BLK_NENTRY, create, &ino)) != 0)
    {
        return ret;
    }
    }

out:
    assert(ino == 0 || sfs_block_inuse(sfs, ino));
    *ino_store = ino;
    return 0;
}
.....
static int
sfs_io_nolock(struct sfs_fs *sfs, struct sfs_inode *sin, void *buf, off_t offset, size_t *alenp, bool write) {
    struct sfs_disk_inode *din = sin->din;
    assert(din->type != SFS_TYPE_DIR);
    off_t endpos = offset + *alenp, blkoff;
    *alenp = 0;
    if (offset < 0 || offset >= SFS_MAX_FILE_SIZE || offset > endpos) {
        return -E_INVALID;
    }
    if (offset == endpos) {
        return 0;
    }
    if (endpos > SFS_MAX_FILE_SIZE) {
        endpos = SFS_MAX_FILE_SIZE;
    }
    if (!write) {
        if (offset >= din->size) {
            return 0;
        }
        if (endpos > din->size) {
            endpos = din->size;
        }
    }

    int (*sfs_buf_op)(struct sfs_fs *sfs, void *buf, size_t len, uint32_t blkno, off_t offset);
    int (*sfs_block_op)(struct sfs_fs *sfs, void *buf, uint32_t blkno, uint32_t nblks);
    if (write) {
        sfs_buf_op = ...(7.1)..., sfs_block_op = ...(7.2)...;
    }
}

```

```

else {
    sfs_buf_op = sfs_rbuf, sfs_block_op = sfs_rblock;
}

int ret = 0;
size_t size, alen = 0;
uint32_t ino;
uint32_t blkno = offset / SFS_BLKSIZE;
uint32_t nblks = endpos / SFS_BLKSIZE - blkno;

if ((blkoff = offset % SFS_BLKSIZE) != 0) {
    size = (nblks != 0) ? (SFS_BLKSIZE - blkoff) : (endpos - offset);
    if ((ret = sfs_bmap_load_nolock(sfs, sin, blkno, &ino)) != 0) {
        goto out;
    }
    if ((ret = ...(8)... ) != 0) {
        goto out;
    }
    alen += size;
    if (nblks == 0) {
        goto out;
    }
    buf += size, blkno ++, nblks --;
}

size = SFS_BLKSIZE;
while (nblks != 0) {
    if ((ret = sfs_bmap_load_nolock(sfs, sin, blkno, &ino)) != 0) {
        goto out;
    }
    if ((ret = sfs_block_op(sfs, buf, ino, 1)) != 0) {
        goto out;
    }
    alen += size, buf += size, blkno ++, nblks --;
}

if ((size = endpos % SFS_BLKSIZE) != 0) {
    if ((ret = sfs_bmap_load_nolock(sfs, sin, blkno, &ino)) != 0) {
        goto out;
    }
    if ((ret = sfs_buf_op(sfs, buf, size, ino, 0)) != 0) {
        goto out;
    }
    alen += size;
}

```

```

out:
    *alenp = alen;

    if (offset + alen > sin->din->size) {
        sin->din->size = offset + alen;
        sin->dirty = 1;
    }

    return ret;
}

.....

static int
sfs_read(struct inode *node, struct iobuf *iob) {
    return ...(9)...;
}

static int
sfs_write(struct inode *node, struct iobuf *iob) {
    return ...(10)...;
}

.....

static inline int
sfs_io(struct inode *node, struct iobuf *iob, bool write) {
    struct sfs_fs *sfs = fsop_info(vop_fs(node), sfs);
    struct sfs_inode *sin = vop_info(node, sfs_inode);
    int ret;
    lock_sin(sin);
    {
        size_t alen = iob->io_resid;
        ret = ...(11)...;
        if (alen != 0) {
            iobuf_skip(iob, alen);
        }
    }
    unlock_sin(sin);
    return ret;
}

.....

static const struct inode_ops sfs_node_dirops = {
    .vop_magic           = VOP_MAGIC,
    .vop_open            = sfs_opendir,
    .vop_close           = sfs_close,
    .vop_read            = NULL_VOP_ISDIR,
    .vop_write           = NULL_VOP_ISDIR,
    .vop_fstat           = sfs_fstat,
    .vop_fsync           = sfs_fsync,
    .vop_mkdir           = NULL_VOP_UNIMP,

```

```

.vop_link                = NULL_VOP_UNIMP,
.vop_rename              = NULL_VOP_UNIMP,
.vop_readlink            = NULL_VOP_ISDIR,
.vop_symlink             = NULL_VOP_UNIMP,
.vop_namefile            = sfs_namefile,
.vop_getdirent          = sfs_getdirent,
.vop_reclaim             = sfs_reclaim,
.vop_ioctl               = NULL_VOP_INVALID,
.vop_gettype             = sfs_gettype,
.vop_tryseek             = NULL_VOP_ISDIR,
.vop_truncate            = NULL_VOP_UNIMP,
.vop_create              = NULL_VOP_UNIMP,
.vop_unlink              = NULL_VOP_UNIMP,
.vop_lookup              = sfs_lookup,
.vop_lookup_parent       = NULL_VOP_UNIMP,
};

static const struct inode_ops sfs_node_fileops = {
    .vop_magic            = VOP_MAGIC,
    .vop_open             = sfs_openfile,
    .vop_close            = sfs_close,
    .vop_read             = sfs_read,
    .vop_write            = sfs_write,
    .vop_fstat            = sfs_fstat,
    .vop_fsync            = sfs_fsync,
    .vop_mkdir            = NULL_VOP_NOTDIR,
    .vop_link             = NULL_VOP_NOTDIR,
    .vop_rename          = NULL_VOP_NOTDIR,
    .vop_readlink         = NULL_VOP_NOTDIR,
    .vop_symlink          = NULL_VOP_NOTDIR,
    .vop_namefile         = NULL_VOP_NOTDIR,
    .vop_getdirent        = NULL_VOP_NOTDIR,
    .vop_reclaim          = sfs_reclaim,
    .vop_ioctl            = NULL_VOP_INVALID,
    .vop_gettype          = sfs_gettype,
    .vop_tryseek          = sfs_tryseek,
    .vop_truncate         = sfs_truncfile,
    .vop_create           = NULL_VOP_NOTDIR,
    .vop_unlink           = NULL_VOP_NOTDIR,
    .vop_lookup           = NULL_VOP_NOTDIR,
    .vop_lookup_parent    = NULL_VOP_NOTDIR,
};

```

六、(15 分)某计算机系统中有 M 个同类型共享资源，有 N 个进程竞争使用，每个进程最多需要 K 个共享资源。该系统不会发生死锁的 K 的最大值是多少？要求给出计算过程，并说明理由。

$$M = N(K-1) + 1$$